

Can we trust this?

A Campus investigation in four 45-minute sessions

Grades 6-8 with a teacher or facilitator

Make an initial judgment. Examine evidence. Choose an action. Explain your reasons. Revisit what changed.

A shared question

How can we make a responsible recommendation when an AI-style reply sounds convincing but the evidence is incomplete?

Four sessions

1. Follow the evidence - Library 2. A better prompt is only a start - Prompt Workshop 3. Make a decision together - Campus discussion 4. Use the method somewhere new - Library and Passport

Choose your format

Paper only: print the student case and response pages. Projected: open site/trust/index.html in a browser; it works offline after downloading. App: open Can we trust this? from Progress, Settings, or a temporary Classroom Mode session. The Library and Prompt Workshop investigations also link to it.

Draft for educator review

Use only with school approval. This packet is a facilitated learning pilot, not a validated curriculum assessment or a claim of district approval. Do not require students to open a live AI service.

All people, organisations, data and AI replies in these cases are fictional and authored for learning. No live AI is used. Do not treat the examples as gardening or product advice.

Copying permission

Teachers may print and copy this packet for their own classroom educational use. Do not resell it or imply school endorsement. Optional paid facilitation is separate from the free learner app and these sample materials.

Prepare once. Discuss together.

Before session 1

Check school approval, device rules and required accommodations. Print the two student pages per session and baseline sheet. Read the teacher notes. Arrange pairs or groups of three. A whiteboard, pencils and printed pages are sufficient.

Screen and paper use

The 45 minutes includes discussion and paper, not 45 minutes of individual screen time. The eight-minute evidence segment can use a projector or printed cards. Follow the school's total daily screen rules. Temporary app sessions clear on background, timeout and ending; keep individual work on paper.

Roles that rotate

Reader: makes the claim and evidence accessible. Checker: asks what a source establishes and what is missing. Reporter: paraphrases others and records the group recommendation. Rotate roles so everyone contributes.

Access and participation

Read cards aloud; offer enlarged copies, oral answers, scribing and familiar assistive tools. Do not score spelling, speed, handwriting, confidence or length. A short supported explanation can be stronger than a long answer. Allow a fictional response rather than personal disclosure.

Individual evidence before group discussion

Keep baseline and session 4 responses individual before feedback. Keep agreed supports comparable and record assistance in aggregate. Keep identified student work at school. Private Passport reflections are optional and are not collected by Vizancia.

What success would look like

Students point to evidence, qualify claims, choose help deliberately and name a useful next check. Teachers can finish a session in the agreed time and decide whether they would use the next sequence. Neither XP nor completion proves learning.

Before we investigate

Fictional case: a supplier says its AI lunch-queue sign halves waiting time for every school. On one Tuesday, one queue using the sign averaged 4 minutes; a different queue on Monday averaged 8. Menus and staffing differed. The trial did not measure how many students got lunch. An AI-style summary repeats the claim without naming a source.

Would you recommend using the sign now? Explain your reasoning, using a detail from the case.

What does this evidence leave uncertain? What would you check next?

Would you use AI, another tool, your own work, or a person to help? Why?

Use the support agreed with your teacher. Leave out names and personal details. Your teacher may keep a class code separately.

1. Follow the evidence

Separate a claim from supporting evidence and identify a limit.

The claim

The Campus newspaper receives this pitch: 'An AI watering system saves 50% of water in every garden. Our two-day test proves the school should buy it.'

Before reading the cards: Would you publish that claim now? Give an initial reason and one thing you would need to know.

A. Groundskeeper's test log

Campus groundskeeper / fictional local record / 4 May

Two similar beds were observed for two days. The usual schedule used 20 litres; the new system used 10. Only water use was recorded. Plant health and longer-term weather were not measured.

B. Supplier's poster

Sprout Systems / fictional advertisement / 5 May

'Save half your water everywhere!' The poster repeats the Campus test. The supplier sells the system. It gives no additional tests or methods.

C. Groundskeeper's clarification

Campus groundskeeper / fictional interview note / 6 May

'The result is promising for these beds. I have not tested a heat wave, different plants or a whole growing season. I would not promise the same savings everywhere.'

Select the evidence cards that matter to your decision. For each one, explain what it can and cannot support.

Session 1: my reasoning

My initial judgment, before the evidence:

Evidence I used, what it supports, and a limit:

Possible actions: Publish the universal claim / Report the small test and its limits / Wait and request a better test

My recommendation and why. Include a useful next check:

Help I chose (own work, person, other tool, AI-style reply, or unsure) and why:

What changed or remains uncertain?

Discuss: Partner A explains the strongest evidence. Partner B asks what is missing. Swap roles. Agree on one sentence that could responsibly appear in the newspaper.

Optional private Passport reflection: What changed between your first judgment and your decision? Name a source and a limit.

2. A better prompt is only a start

Distinguish a useful response from a verified claim and choose appropriate help.

The claim

The newspaper team wants a recommendation about the same watering system. Two authored AI-style replies use the Library evidence. Neither reply is a new source.

Before reading the cards: Before reading the replies, write two things a responsible recommendation must include. Would you use AI, another tool, a person, or your own work for this task?

A. Reply to 'Make the system sound exciting'

Authored simulation / not a live model or evidence

'The AI system guarantees a 50% saving in every garden. Its successful test proves it is ready for the whole school. Buy it now!'

B. Reply to 'Separate observations, limits and next steps'

Authored simulation / not a live model or evidence

'The two-day test used 10 rather than 20 litres. We do not know the result in other conditions. Run a longer comparison and measure plant health. The system costs 30 credits.'

C. Source check

Campus evidence folder / fictional document inventory

The folder contains the two-day log, supplier poster and groundskeeper clarification. None gives a price. The 30-credit price in reply B cannot be verified from these documents.

Compare the replies with the test log. Select the cards you used to check a claim. Notice what a clearer prompt improved and what it could not verify.

Session 2: my reasoning

My initial judgment, before the evidence:

Evidence I used, what it supports, and a limit:

Possible actions: Use reply B unchanged / Check B, remove the price, and write our conclusion / Work from the records with a person instead

My recommendation and why. Include a useful next check:

Help I chose (own work, person, other tool, AI-style reply, or unsure) and why:

What changed or remains uncertain?

Discuss: One partner checks each factual sentence; the other proposes wording. Swap. Discuss a task where AI help would be useful and a task where it would skip the skill you need to practise. Follow the school's rules.

Optional private Passport reflection: Which help did you choose and why? What did you personally check or contribute?

3. Make a decision together

Consider affected people and propose a proportionate, reversible decision.

The claim

At a Campus meeting, the newspaper must recommend buying the system, running a limited trial, or waiting. There is no deadline requiring a purchase today.

Before reading the cards: Who should be heard before a decision? Write your recommendation and one question you would ask them.

A. Groundskeeper's priorities

Fictional role card / read aloud or discuss

'I want to conserve water and keep plants healthy. I can supervise a small trial. We need a manual stop and a clear rule for checking dry soil. A person must remain responsible.'

B. Student garden club

Fictional role card / read aloud or discuss

'We want to keep learning how to care for the garden. The controls should work for everyone. We do not agree to upload our names or photos to test the watering system.'

C. School purchasing lead

Fictional role card / read aloud or discuss

'No price has been confirmed. We need the real costs, approval and a comparison with a simpler timer. We can pause the proposal while gathering this information.'

Select the perspectives that affect your recommendation. Separate what people value from claims that require testing.

Session 3: my reasoning

My initial judgment, before the evidence:

Evidence I used, what it supports, and a limit:

Possible actions: Buy for every garden immediately / Propose a supervised trial with conditions / Wait or use a simpler solution

My recommendation and why. Include a useful next check:

Help I chose (own work, person, other tool, AI-style reply, or unsure) and why:

What changed or remains uncertain?

Discuss: Use the three role cards. Each person speaks, another paraphrases, and the group records one disagreement fairly. Write a recommendation with a responsible person and a review or stop condition.

Optional private Passport reflection: Whose perspective changed your thinking? What does your group remain responsible for? You may keep personal reflections private and share only your project recommendation.

4. Use the method somewhere new

Apply source checking, limits and responsible tool choice to an unfamiliar claim.

The claim

A fictional company says its AI quiet-desk display doubles reading comprehension for every student. It cites a lunchtime trial and an enthusiastic AI-style summary. The library is considering a purchase.

Before reading the cards: Before opening the evidence, say what you would need to know and how you would check it. Complete your initial response without using the earlier worked answers; keep agreed accessibility supports.

A. Trial note

Fictional library log / 12 June

Six volunteers sat at the display for two lunch breaks. They read twice as many pages as a different group the previous week. No comprehension questions were asked. Reading materials and noise levels differed.

B. AI-style summary

Authored simulation / not independent evidence

'The trial confirms doubled comprehension for all students. Several websites agree.' The summary gives no website titles, links or new measurements.

C. Librarian's note

Fictional interview note / 13 June

'Pages read and understanding are different measures. We need a fair comparison and accessible options. Some students prefer a quiet corner without a display. We have not obtained a price or approval.'

Make a final recommendation using at least one source, one limitation and a next check. Decide what help is appropriate. Explain aloud, on paper or in your Passport.

Session 4: my reasoning

My initial judgment, before the evidence:

Evidence I used, what it supports, and a limit:

Possible actions: Recommend buying on this evidence / Request a fair test and involve library users / Decline for now and explain what would change my mind

My recommendation and why. Include a useful next check:

Help I chose (own work, person, other tool, AI-style reply, or unsure) and why:

What changed or remains uncertain?

Discuss: Give your individual recommendation before viewing the feedback. Then compare with a partner: what reasoning carried over from the garden case, and what was different? Teachers may score the assigned response using the paper rubric.

Optional private Passport reflection: How did you decide what to trust? What can you now explain without copying the previous example? What will you check in a real task?

Assess the explanation

Use this informal rubric on the assigned response before showing feedback. Score each dimension 0, 1 or 2. A total of 0-8 is a description of this task, not a grade of a person or an AI-literacy certification.

Evidence

0: no relevant evidence or unsupported claim. 1: refers to a relevant detail. 2: explains what that source actually supports and distinguishes repetition from independent support.

Limits

0: accepts or rejects without a reason. 1: names a relevant uncertainty. 2: explains why a measurement, comparison, sample or duration limits the conclusion.

Next action

0: no useful next step. 1: suggests a relevant check. 2: specifies a feasible check and connects it to the decision or people affected.

Help and responsibility

0: no reason for tool choice. 1: gives a relevant reason for using or declining help. 2: connects that choice to permission, learning needs or verification and keeps a person responsible.

Interpretation

Score oral and written responses on the same reasoning criteria. Do not score agreement with a preferred option. A justified pause, qualified report or proposed test may all be defensible. Keep individual results at school; report group totals only when safe.

Avoid overstating change

The cases are related but not equivalent standardised tests. Practice, teacher help, attendance and differing tasks can affect scores. Pre/post change without a comparison group does not establish that Vizancia caused it. Missing responses are not zeros.

Facilitator notes: sessions 1 and 2

Session 1: Follow the evidence

A supports a narrow description of observed water use, not plant health or universal savings. B is a sales claim based on A, not independent corroboration. C narrows the interpretation. Accept a qualified report or a justified pause; score reasoning, not agreement with one preferred action.

Publish the universal claim: The figures describe two beds for two days. Repeating them in an advertisement adds no independent evidence. A universal claim goes beyond this test.

Report the small test and its limits: This is supported if you explain the conditions and missing measures. You can report 10 versus 20 litres in this test without promising universal savings or healthy plants.

Wait and request a better test: This is reasonable if the article's purpose is to recommend a purchase. Specify what you need: more conditions, longer observation, plant health and an independent check.

Pacing: 5 min: initial judgment / 8 min: read or project evidence / 12 min: pairs examine sources / 12 min: draft a responsible headline / 8 min: reflect and exit

Session 2: A better prompt is only a start

Reply A overgeneralises and pushes a purchase. Reply B correctly narrows the water result but fabricates a price. Checking is still necessary after prompt improvement. Accept justified use, non-use or human help; do not award extra credit for using AI. Accessibility support is not cheating.

Use reply B unchanged: B is more useful than A, but it adds an unsupported price. Structure and caution do not guarantee that every detail is grounded.

Check B, remove the price, and write our conclusion: This keeps useful organisation while correcting the unsupported detail. If AI help is permitted, disclose its role; your group remains responsible for the final recommendation.

Work from the records with a person instead: This is a valid choice. A librarian or groundskeeper can help examine the original records. Explain how you will still check facts and avoid inventing missing information.

Pacing: 5 min: set criteria before replies / 8 min: compare authored replies / 12 min: verify each factual sentence / 12 min: write a checked recommendation / 8 min: reflect on help and responsibility

Facilitator notes: sessions 3 and 4

Session 3: Make a decision together

Look for listening, a distinction between evidence and values, privacy, accessibility, a comparison and a reversible decision. More than one recommendation is defensible. Never score a student's beliefs or require private disclosures. The teacher evaluates the assigned recommendation only.

Buy for every garden immediately: The evidence does not establish cost, accessibility or performance across gardens. This commits resources before the unanswered questions are resolved.

Propose a supervised trial with conditions: Name the responsible person, required approval, comparison, success measures, stop rule and review date. A proposal is not permission to start the trial.

Wait or use a simpler solution: A justified pause or simpler timer is reasonable. Explain what need it meets, what is still unknown, and what evidence could change your mind.

Pacing: 5 min: initial recommendation / 8 min: read role cards / 12 min: meeting with rotating roles / 12 min: write a conditional recommendation / 8 min: reflect and identify a next check

Session 4: Use the method somewhere new

Score the initial individual recommendation before feedback. Strong responses distinguish pages from comprehension, identify unequal groups and short duration, reject unsupported corroboration, include affected people and propose a concrete next check. Score the reasoning, not the selected option. This is an informal transfer check, not a validated assessment or proof of programme effectiveness.

Recommend buying on this evidence: The trial did not measure comprehension and did not compare equivalent conditions. Unnamed websites and a fluent summary cannot establish the claim.

Request a fair test and involve library users: Specify an understanding measure, comparable conditions, accessible participation, permission and an independent check. A proposed test could inform a later decision; it does not guarantee a benefit.

Decline for now and explain what would change my mind: This can be justified by the missing evidence, price or unmet needs. Explain the limits precisely and name what new evidence would be relevant.

Pacing: 5 min: individual first judgment / 8 min: read new evidence / 12 min: individual transfer response before feedback / 12 min: partner explanation and feedback / 8 min: Passport revisit plan

Pilot observation and next steps

Pilot code, session number, date and delivery format:

Preparation minutes / teaching minutes / interruptions / accommodations used (aggregate only):

How many attended? How many completed an individual explanation? Which step was unclear?

What would you change? Would you independently schedule another session? Why?

After the sequence

Review aggregate rubric counts, missing responses, teacher preparation, founder support time and a concrete continuation decision. Record willingness to use separately from willingness to pay. A booked continuation or paid invoice is stronger evidence than positive comments.

No participant has been recruited or approved by creation of these materials. Teacher interest, school approval, scheduling and payment are separate steps.